

Vantari — Round 3 Simulation Plan

Oct 2, 2026 · @Eshwar

What Wizard of Oz means here

Wizard of Oz is a method from early human-computer interaction research. A person interacts with what they believe is a working system, while a hidden operator supplies the parts that do not exist yet. The point is to test the experience before building the technology.

The Ken has inverted it. Normally the system is faked and the human hides. Here the agent must genuinely run, and the world around it is played.

What runs for real. The model, with our system prompt, making every decision itself. Gnani, called properly, with whatever it returns pasted back untouched. WhatsApp, connected. A Google Sheet holding what the agent knows.

What is played behind the curtain. Pine Labs and Delhivery, answered by a teammate using the response shape their documentation specifies. The clinic receptionist. The chemist. Every inbound event from the outside world.

Three rules that decide whether this passes.

1. The person behind the curtain never helps the agent think. The moment a teammate types "now book the appointment", it stops being an agent.
2. Every input must have a real source. A WhatsApp voice note qualifies. A clock reaching 9am qualifies. A physical SOS button does not exist, so it cannot produce an input.
3. We get three imagined capabilities, no more. Round 2 listed considerably more than three must-builds, so the list has to be cut.

The scenario we film

Lakshmi, 67, lives alone in Warangal. Her son Arun, 29, works in Hyderabad. She has hypertension and takes two tablets a day. Her last BP check was four months ago.

On a Tuesday morning she sends a voice note in Telugu to the family WhatsApp group saying she felt giddy when she stood up, and that her tablets are nearly finished. Everything that follows runs from those two facts.

Why this episode and not another.

It matches the loop we submitted in Round 2, which triggers on a health issue raised by the patient or the responsible person. It carries all three rails in one run: a call to a clinic, a payment at a chemist, a delivery to her door. And it closes with a doctor appointment confirmed and medicine dispatched, which is an outcome a viewer can see arrive rather than take on trust.

It also lets the refill clock appear without being the trigger. She mentions the tablets herself, so the agent learns the dispensed quantity and sets the next cycle. That is our strongest evidence and our headline rail gap, and it would be lost if we filmed a pure symptom episode.


What we deliberately leave out. The cab, because no rail covers moving a person. The SOS button, because no real source produces that input today. Mid-transit redirect, because it costs screen time and carries nothing.

The alternative, if five minutes proves too tight. Film only the refill cycle, triggered by the clock. It is cleaner and shorter, but it loses the voice rail almost entirely, which is the part the jury will want to see working.

The flow, beat by beat

One voice note ends with an appointment booked and medicine dispatched

Tuesday, 9:12am. Lakshmi in Warangal, Arun in Hyderabad.



the filmed episode · 9 steps, 3 rails

The loop closes on itself. The last step sets a clock that triggers the next cycle without anyone sending anything.

0:00 to 0:40 — the trigger. Lakshmi records a voice note in Telugu to the family group. On screen: the note arriving in WhatsApp. The agent passes the audio to Gnani and the transcript comes back. We show the raw Gnani response, untouched.

0:40 to 1:20 — triage. The agent reads the transcript against what it already holds: hypertension, two tablets daily, last BP check four months ago, treating doctor in Warangal. It decides this is urgent but not an emergency, and that two things are needed, a consultation and a refill. It does not translate her words into clinical terms, and the transcript it carries forward is hers.

1:20 to 2:10 — the call. The agent calls the clinic. A teammate plays the receptionist. The agent speaks through Gnani text-to-speech, says who it is calling for, requests a slot with her existing doctor, waits through hold. When a human answers it rings Lakshmi and hands her the line. She confirms the time herself, in Telugu.

2:10 to 2:50 — the money. The chemist bill is above the standing limit. The agent messages Arun on WhatsApp naming the medicine and the amount. He replies yes. The agent pays through Pine Labs and receives back the drug, strength and quantity actually dispensed.

2:50 to 3:30 — the delivery. Serviceability checked for her pincode, waybill generated, shipment created, tracked to delivered, with confirmation that Lakshmi herself received it.

3:30 to 4:10 — the family sees it. Every action posts to the group as it happens. Arun's sister, who did nothing, asks the agent in the group when the last BP check was. It answers. This is our Round 2 insight on screen.

4:10 to 4:40 — the loop closes. From the quantity dispensed, not the quantity prescribed, the agent calculates the date the tablets run out and sets the next cycle. Nobody has to remember anything.

Component checklist

Y runs for real on camera. **N** played behind the curtain, using the documented response. **B** one of our three imagined capabilities.

The agent itself

#	Component	Y/N/B	Note
1.1	The model making every decision	Y	Our system prompt. Nobody tells it the next step.
1.2	System prompt holding the rules	Y	Limits, escalation order, the clinical-term guardrail, who gets asked when.
1.3	Family profile on day one	Y	Conditions, medicines, doctor, contacts, spend limit. Seeded by Arun at signup.
1.4	Memory of the last cycle	Y	A Google Sheet. Real connector, visible on screen.

Trigger

#	Component	Y/N/B	Note
2.1	WhatsApp voice note from Lakshmi	Y	Real source, real arrival. Our opening shot.
2.2	Physical SOS button	N	No real source produces it today. Cut from the recording.
2.3	Clock reaching a refill date	Y	Real clock. Closes the loop at the end.
2.4	Message from Arun	Y	WhatsApp, real.

Voice rail, Gnani

#	Component	Y/N/B	Note
3.1	Speech to text on the voice note	Y	Our key covers it. Paste back exactly what returns.
3.2	Text to speech for the call and readback	Y	Same key.

#	Component	Y/N/B	Note
3.3	Telugu and code-switched transcription	Y	Record the note in Telugu. Demonstrates the rail and the insight at once.
3.4	Placing the call, hold detection, warm transfer	B	Imagined capability one. Our key is speech only, not telephony.
3.5	Clinical-term guardrail	Y	System prompt, not a rail. Costs nothing.

Payments rail, Pine Labs

#	Component	Y/N/B	Note
4.1	Grant token with scopes and limits	N	Documented. Teammate returns the documented shape.
4.2	Payment on a 402 challenge, receipt back	N	Documented.
4.3	Agent stopping above the limit and asking	Y	Our prompt enforcing it, visible on screen.
4.4	Itemised dispensing receipt	B	Imagined capability two. Our headline rail gap. Do not drop.
4.5	Cross-person mandate, Arun pays for Lakshmi	N	Argue it is a configuration of the documented mandate, not a new endpoint. Saves a slot.

Logistics rail, Delhivery

#	Component	Y/N/B	Note
5.1	Pincode serviceability	N	Documented. Returns serviceable, COD and prepaid.
5.2	Waybill and order creation	N	Documented.
5.3	Tracking to delivered	N	Documented, pull and webhook.

#	Component	Y/N/B	Note
5.4	Named-recipient confirmation	B	Imagined capability three. Proves the medicine reached Lakshmi, not the gate.
5.5	Mid-transit redirect	N	Cut. Costs screen time, carries nothing.
5.6	Cab for the patient	N	No rail covers moving a person. Out of the recording.

Records and interface

#	Component	Y/N/B	Note
6.1	Reports already in the WhatsApp group	Y	Real, and how families actually do it.
6.2	History file assembled for the doctor	Y	Agent composes it from what it holds. No rail needed.
6.3	Pulling history from ABDM or a hospital	N	Our fourth-rail answer. Flag the gap in the write-up, do not spend a slot.
7.1	WhatsApp to Arun	Y	Connected. Real messages on screen.
7.2	Group posting every action	Y	Real. Our Round 2 insight made visible.
7.3	Family asking the agent questions in the group	Y	Real. Worth twenty seconds.
7.4	Clinic receptionist	N	Teammate plays her. Agent speaks through Gnani.
7.5	Approval buttons	N	Plain WhatsApp replies. Do not fake a UI.

The three imagined capabilities

The brief allows three. Each needs a partner, an endpoint and a statement of what that partner already holds that makes it possible.

One. Agent-placed call with hold detection and warm transfer. Gnani.

Endpoint shape: a call session that takes a destination number, streams the agent's synthesised speech, reports when a human has answered, and transfers the leg to a second number. What Gnani already holds: production speech recognition and synthesis tuned for 8kHz telephony across ten Indian languages, already deployed in enterprise call automation. The models work on live calls today. What is missing is agent-initiated call control, dialling out and handing the line over.

Why this one matters: twelve of the twenty-five people we surveyed said their parent no longer phones a clinic. The call is the exact errand that needs automating, and it is the piece the rail does not reach.

Two. Itemised dispensing receipt. Pine Labs.

Endpoint shape: the payment receipt returns line items, drug name, strength and quantity dispensed, alongside the amount. What Pine Labs already holds: the merchant relationship with the pharmacy and the authorisation record for every transaction at that counter. The data passes through them. It is simply not returned.

Why this one matters: the doctor we interviewed described patients carrying five hundred rupees and buying whatever that covers of a thirty-day prescription. If the refill clock is computed from the prescription it is wrong. Only the counter knows the real number.

Three. Named-recipient confirmation. Delhivery.

Endpoint shape: the delivery event returns which named person took the parcel, against a list the sender supplied. What Delhivery already holds: OTP-based delivery confirmation and a complete event trail per waybill. The identity is established at the door. It is not returned to the party who ordered.

Why this one matters: our agent closes an episode on medicine reaching Lakshmi, not on a parcel reaching an address. Prescription medicine left with a neighbour is not a completed episode.

What we deliberately did not spend a slot on. The cross-person mandate, argued instead as a configuration of the documented Pine Labs mandate. Mid-transit redirect. The health records rail, which we name as a fourth rail in the write-up rather than imagining an endpoint for it.

Before filming

To build.

- The system prompt. Every rule the agent follows has to be written down, because Part 1 asks us to name the rule behind each decision. Write it before anything else.
- The family profile sheet. Conditions, medicines with quantities, doctor, contacts, spend limit. This is the day-one answer: Arun enters it at signup.
- Gnani API key, tested on a real Telugu recording before filming day.
- WhatsApp connected, with a real group containing at least three people.
- A decision log the agent writes to as it goes. Part 1 wants every decision with a timestamp, the input, the source, the rule and the exact message sent. Capturing it live beats reconstructing it afterwards.

To record.

- The voice note, in Telugu, under fifteen seconds. Giddiness on standing, tablets nearly over.
- Teammate scripts for the receptionist and the chemist. They respond as those people would, and never as helpers.
- Documented response shapes for Pine Labs and Delhivery, copied out in advance so nothing is invented live.

Open decisions.

1. Whether the sister's question in the group earns its forty seconds. It carries our strongest insight but it is not on the critical path to the outcome.
2. Whether to show the agent failing once. A no-slot or out-of-stock branch would demonstrate judgement, but five minutes is tight.
3. Whether Lakshmi confirms the slot on the call or afterwards. On the call is truer to the design and costs more seconds.

Honest risks.

The five-minute limit is the binding constraint, not the design. Three rails, two humans and nine steps is a lot to show. If it runs long, the delivery leg compresses best, since tracking is the least surprising part.

The partial-dispensing claim behind imagined capability two still rests on one clinician. One phone call to a chemist asking whether a part-fill is itemised on the bill either hardens it or tells us to reframe it.